Member-only story

The 9 Most Critical Plots Every Data Scientist Must Master

Roshmita Dey Follow 5 min read · 4 days ago

96

If you strip data science down to its core, it's really about **making decisions under uncertainty**. Plots are the fastest way to reduce that uncertainty — whether you're diagnosing a model, explaining results to stakeholders, or choosing between competing approaches.

Below are **nine essential plots** every data scientist should deeply understand — not just how to generate them, but **how to interpret and use them correctly in practice**.

1. KS Plot (Kolmogorov–Smirnov Plot)

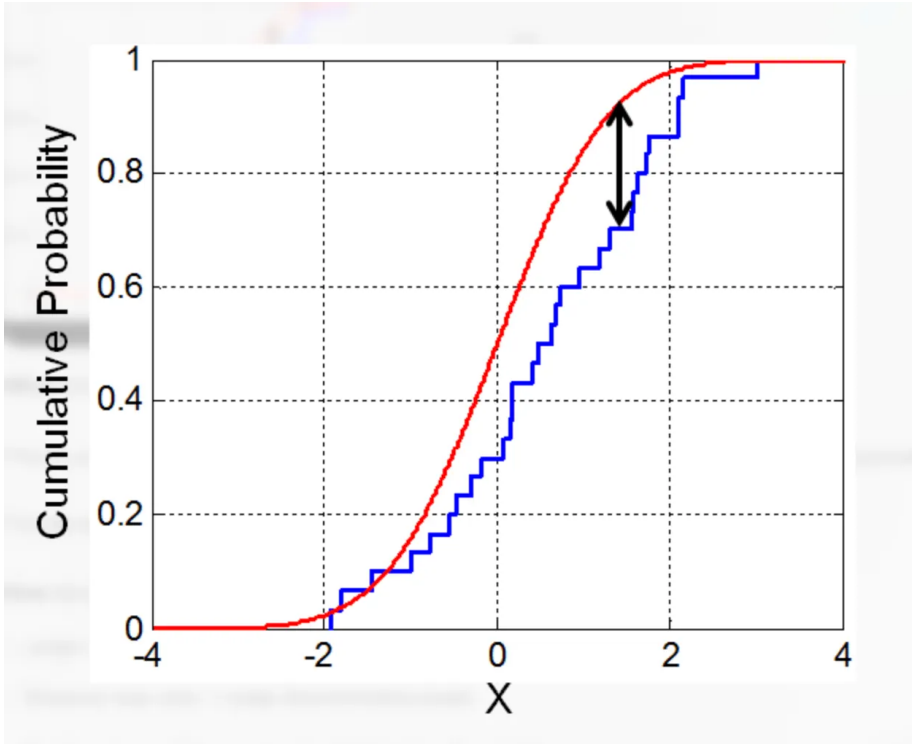

What it is

A KS plot compares the **cumulative distribution functions (CDFs)** of prediction scores for two classes (usually positive vs negative).

The **KS statistic** is the **maximum vertical distance** between the two curves.

How to interpret

- Larger KS distance → better separation between classes
- KS close to 0 → model barely distinguishes positives from negatives
- The x-value at max separation often suggests a **good threshold**

When to use

- Credit risk, fraud detection, churn models
- Highly popular in **banking and risk modeling**
- When you care about **rank ordering quality**, not just accuracy

> **Rule of thumb:**
>
> *KS < 20 → weak*
>
> *20–40 → decent*
>
> *40 → strong (domain dependent)*

2. SHAP Plot

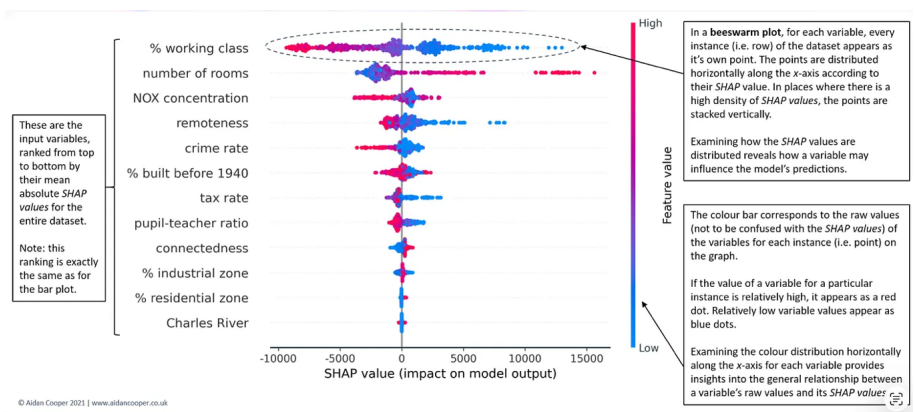

What it is

SHAP plots show **feature contributions** based on cooperative game theory. They explain **how much each feature pushes a prediction up or down.**

Common variants:

- Bar plot → global importance

- Beeswarm → distribution + direction of impact

How to interpret

- Positive SHAP value → feature increases prediction
- Negative SHAP value → feature decreases prediction
- Wider spread → non-linear or interaction effects

When to use

- Model explainability (especially for tree-based models)
- Regulatory or high-stakes models
- Debugging "why is the model behaving weirdly?"

> *Big win: SHAP explains individual predictions, not just averages.*

3. Q–Q Plot (Quantile–Quantile Plot)

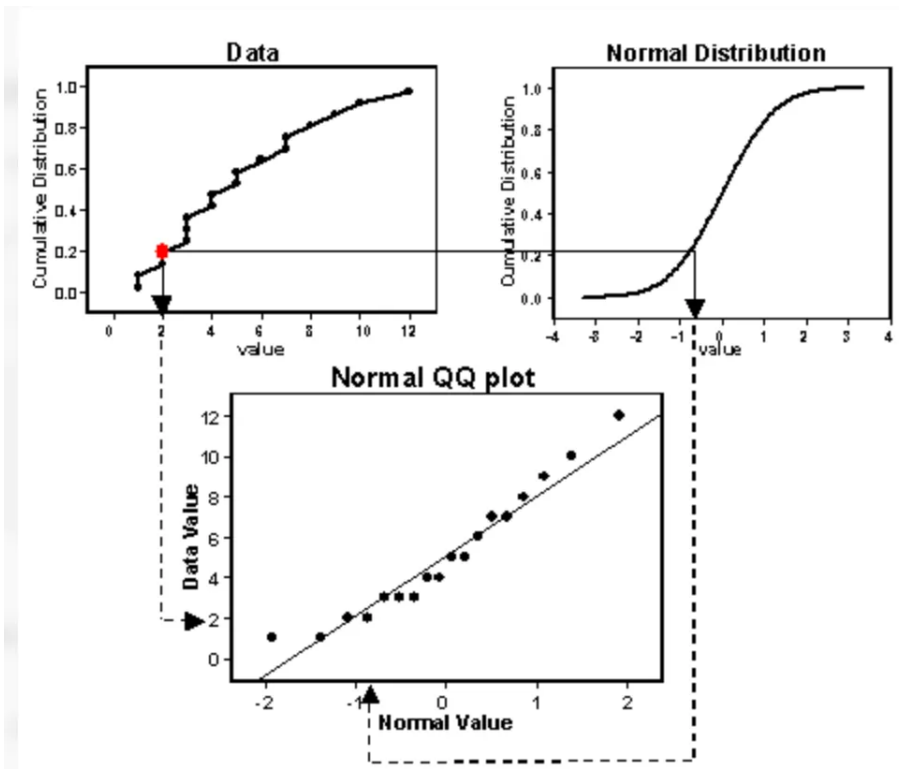

What it is

A Q–Q plot compares **sample quantiles** to **theoretical quantiles** (usually normal distribution).

How to interpret

- Points on straight line → data follows assumed distribution
- S-shape → skewness

- Heavy tails → outliers or non-normality

When to use

- Before applying parametric tests
- Checking regression residual assumptions
- Validating normality assumptions in statistics

> *Pro tip: Normality of residuals, not features, usually matters more.*

4. Cumulative Explained Variance Plot

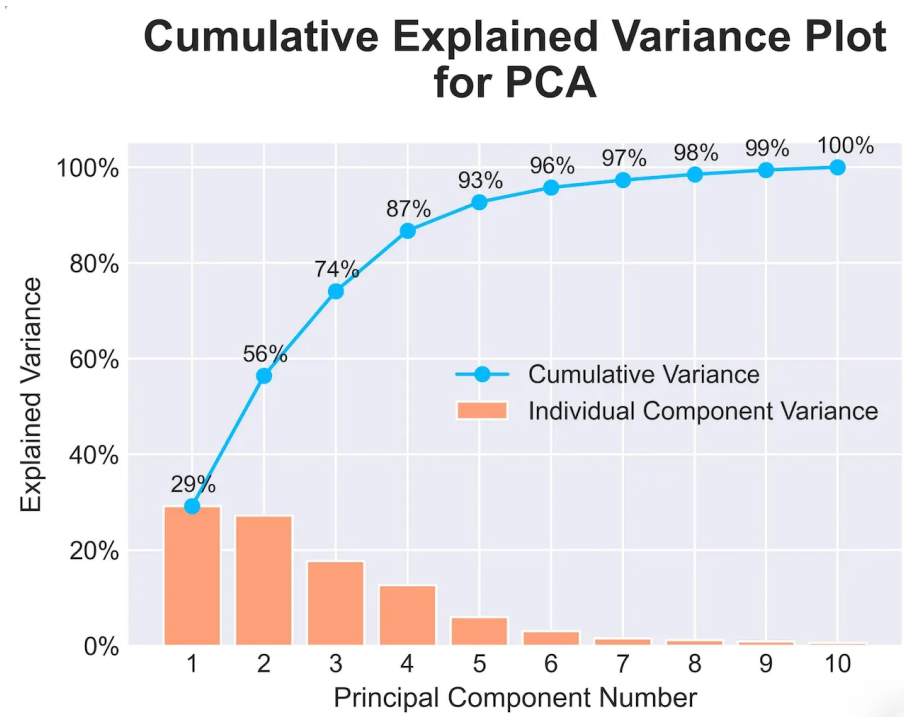

What it is

Used in PCA to show how much **variance is explained** as you add more components.

How to interpret

- Steep rise → few components capture most information
- Flat curve → diminishing returns

When to use

- Dimensionality reduction
- Feature compression
- Noise removal

> *Common practice: Choose components explaining 90–95% variance, not 100%.*

5. Gini Impurity vs Entropy

What it is

These are **split quality metrics** used in decision trees.

- **Gini impurity**: probability of misclassification
- **Entropy**: information disorder

How to interpret

- Both are lowest when node is pure
- Entropy penalizes uncertainty more strongly

When to use

- Decision tree modeling
- Understanding how tree algorithms choose splits

> *In practice: Performance difference is usually negligible — interpretability matters more.*

6. Bias–Variance Tradeoff Plot

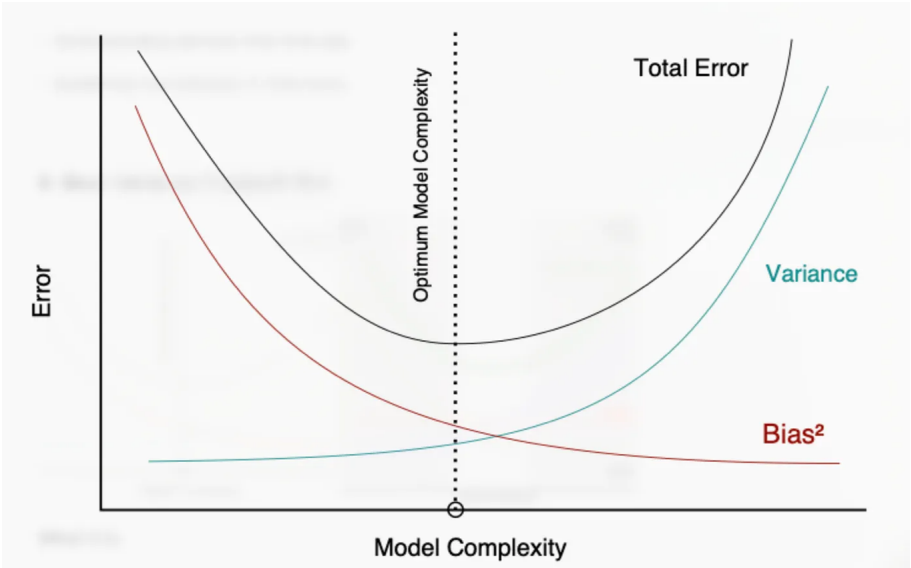

What it is

Illustrates how **model complexity** affects:

- Bias² (underfitting)
- Variance (overfitting)
- Total error

How to interpret

- Left side → high bias, low variance
- Right side → low bias, high variance
- Sweet spot → minimum total error

When to use

- Model selection
- Explaining overfitting to non-technical stakeholders
- Deciding between simple vs complex models

> *Real-world insight: Regularization shifts the curve left or right.*

7. ROC Curve

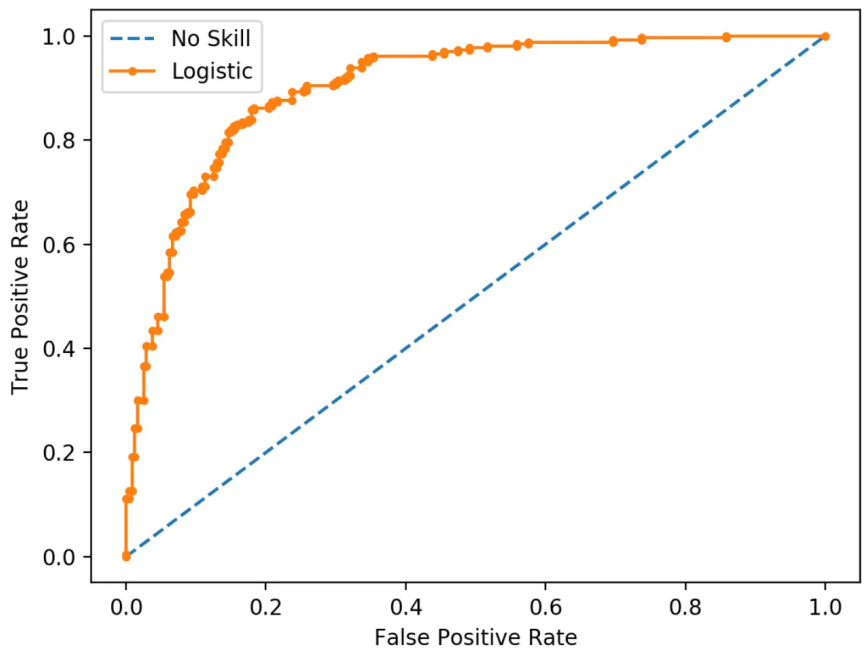

What it is

Plots **True Positive Rate vs False Positive Rate** across thresholds.

How to interpret

- Curve closer to top-left → better
- Diagonal line → random guessing
- AUC measures overall ranking ability

When to use

- Balanced classification problems
- Comparing models independent of threshold

> *Caution: ROC can look overly optimistic on imbalanced datasets.*

8. Precision–Recall Curve

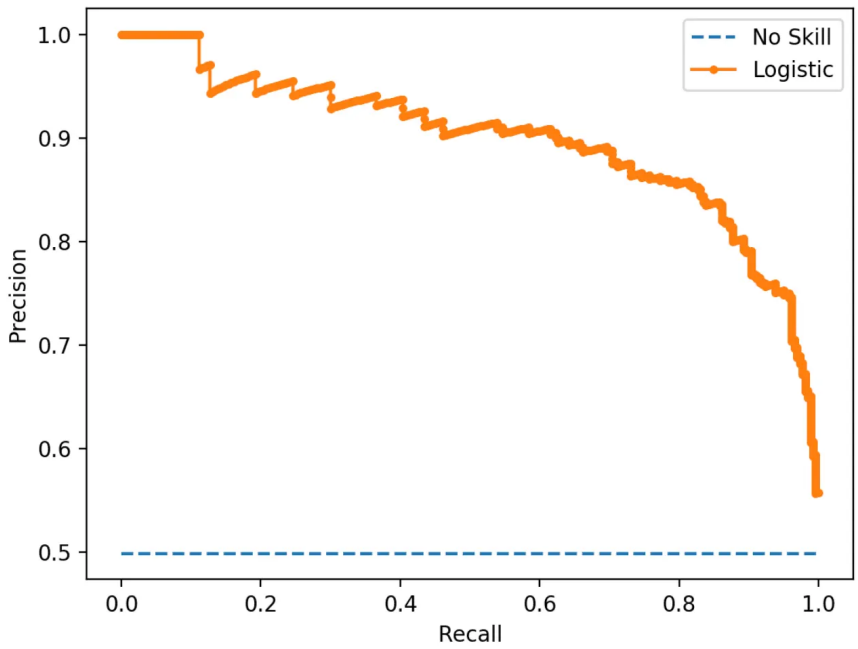

What it is

Plots **Precision vs Recall** across thresholds.

How to interpret

- High curve → good positive-class performance
- Steep drop → precision suffers when recall increases

When to use

- Imbalanced datasets
- Fraud, medical diagnosis, anomaly detection

> *Key insight: If positives are rare, PR curve > ROC curve.*

9. Elbow Curve

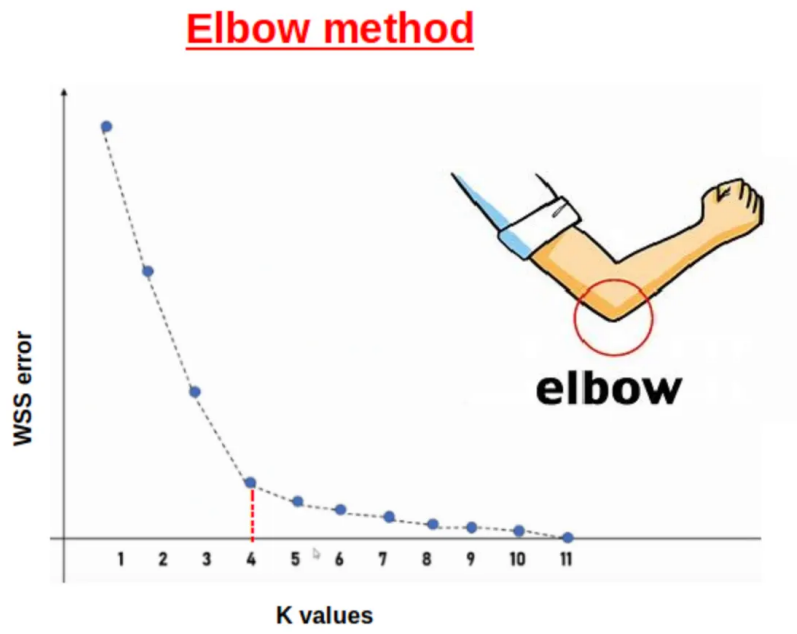

What it is

Plots **within-cluster sum of squares (WCSS)** against number of clusters.

How to interpret

- Sharp bend ("elbow") → optimal number of clusters
- No clear elbow → data may not be clusterable

When to use

- K-Means clustering
- Early exploratory data analysis

> *Pair it with: Silhouette score or domain knowledge.*

These plots aren't just academic — they're **decision tools**.

A strong data scientist doesn't just generate them, but:

Medium Search Write 1

Home
Library
Profile
Stories
Stats

- Understands **what can go wrong**
- Can **explain them clearly** to business stakeholders

If you truly master these nine, you'll:

- Debug models faster
- Make better tradeoffs

- Sound confident in interviews and reviews

And most importantly — you'll trust your models for the **righ**

Data Visualization | Machine Learning | Data Science | Machine Learning Tools

Written by Roshmita Dey
539 followers · 37 following
Follow

Working as a Data Scientist in one of the leading Global banks, my expertise is in the field of Statistics and proficiency in Python, SQL and PySpark

No responses yet

Rob J Goedman

What are your thoughts?

More from Roshmita Dey

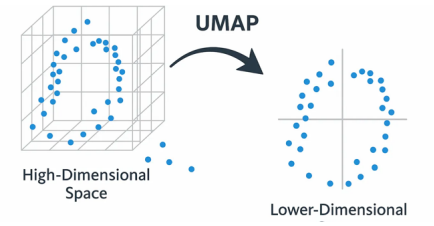

Roshmita Dey

UMAP: The Fast, Powerful Technique Behind Modern...

As datasets grow in size and complexity, visualizing and understanding high-...

Nov 15, 2025 29 1

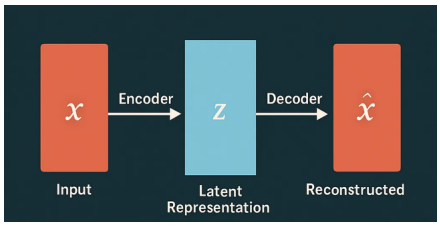

Roshmita Dey

Autoencoders for Anomaly Detection

An autoencoder is a type of neural network designed to learn an identity...

Dec 7, 2025 23

Roshmita Dey

Handling Skewed Data—Make Your Distributions Work for You

Have you ever built a regression or classification model and found your...

Oct 29, 2025 2

Roshmita Dey

Understanding Principal Component Analysis (PCA)

Principal Component Analysis, or PCA, is a fundamental technique in the real...

Oct 6, 2023 224 4

See all from Roshmita Dey

Recommended from Medium

In The Pythonwo... by Aashish Kum...

Stop Writing Utility Functions in Python (Here's the Better Pattern)

Utility functions feel productive—until they quietly turn your clean codebase...

6d ago 356 7

In Write A Cata... by Dr. Patricia Sch...

As a Neuroscientist, I Quit These 5 Morning Habits That Destroy You...

Most people do #1 within 10 minutes of waking (and it sabotages your entire...

Jan 14 20K 346

In Let's Code ... by The Unwritten Al...

Git Confused Me for Years Until I Found This Simple Guide

Most developers don't really understand Git—here's the simple truth.

Dec 18, 2025 2.9K 63

In AWS in Plain En... by Khushbu S...

Zero ETL Is the Reality Check Every Data Engineer Needs in 2026

Why the best data engineers are becoming platform designers?

Jan 19 145 1

MOHAMMED MOSEENA

My ML Interview Notebook: A Revision Guide for Data Science...

ML Revision Notes

Nov 18, 2025 302 11

Havrive Anıl

Exploratory Analysis for Time Series

Hello and welcome to the second article in the Time Series Analysis and...

Jan 21 160

See more recommendations

Help Status About Careers Press Blog Privacy Rules Terms
Text to speech